

# Fundamentals of Computer Vision

Unit 6: Feature Extraction

Jorge Bernal

# Index

01

1. Introduction

02

2. Corner  
detectors

03

3. Edge  
detectors

04

4. Blob  
detectors

1

# Introduction

# Introduction

## Definition of feature:

- Piece of information that is useful to solve a given task
- Interesting part of the image


## Types of features:

- **Global:** global properties of the whole image
  - Mean grey level, mean colour, main colours, histogram
- **Local:** properties of a part of the image with their own entity
  - Points, edges, regions

# Introduction

GLOBAL

LOCAL



# Introduction

- **Local features:**
  - Part of an image that differs from its surroundings.
  - They are associated to a change in a certain property (intensity, color, texture)
  - Examples:
    - Points (corners, interest points)
    - Edges, ridges
    - Small regions (blobs)

# Introduction

- How can we find local features?
  - Feature detection/extraction algorithms
- Detection/Extraction: locate the position of the feature
- Description (Unit 7): measures that are taken from the detected feature that allow us to distinguish it or compare with others

# Introduction

- Why do we use features?
  - They have been used with success in several disciplines and applications:
    - Edge detection associated to roads in aerial images
    - Quality control
    - Polyp Detection
  - Interest points play a key role for certain Applications:
    - Tracking
    - 3D reconstruction
  - They are a first step to achieve a robust image representation:
    - Object recognition
    - Scene classification
    - Texture analysis
    - Image search



# Introduction

- Ideal properties:
  - Repetability:
    - Invariance to transformations
    - Robustness
  - Differentiation (highly different from another)
  - Precise localization
  - Enough points for the needed task
  - Efficient
- Scale: very important factor to achieve robustness, invariance and precision. Allows us to work with different images at several distances.


# Introduction

Feature Detector				Rotation invariant	Scale invariant	Affine invariant	Localization			
	Corner	Blob	Region				Repeatability	accuracy	Robustness	Efficiency
Harris	✓			✓			+++	+++	+++	++
Hessian		✓		✓			++	++	++	+
SUSAN	✓			✓			++	++	++	+++
Harris-Laplace	✓	(✓)		✓	✓		+++	+++	++	+
Hessian-Laplace	(✓)	✓		✓	✓		+++	+++	+++	+
DoG	(✓)	✓		✓	✓		++	++	++	++
SURF	(✓)	✓		✓	✓		++	++	++	+++
Harris-Affine	✓	(✓)		✓	✓	✓	+++	+++	++	++
Hessian-Affine	(✓)	✓		✓	✓	✓	+++	+++	+++	++
Salient Regions	(✓)	✓		✓	✓	(✓)	+	+	++	+
Edge-based	✓			✓	✓	✓	+++	+++	+	+
MSER			✓	✓	✓	✓	+++	+++	++	+++
Intensity-based			✓	✓	✓	✓	++	++	++	++
Superpixels			✓	✓	(✓)	(✓)	+	+	+	+

2


# Corner Detection

# Corner Detection




## Plain region

No changes in all directions



## Edge

No change in edge direction



## Corner

Significant changes in all directions

# Corner Detection

- Harris (1988): Based on the analysis of the 2D structural tensor (second derivative matrix, second moment matrix)
- SUSAN (Smallest Univalued Segment Assimilating Nucleus): morphologic focus
- Harris-Laplace: Use of Harris for a first detection; scale is fixed using laplacian
- Harris-Affine: Use of Harris-Laplace; then it tries to estimate the most affine shape (with an ellipse that is later normalized to a circle)


# Corner Detection: Harris

- In an image intensity corner, intensity changes significantly in all directions.
- Here we are focused in intensity changes in a local window.
- We use SSD: sum of squared differences

$$S(x, y) = \sum_u \sum_v w(u, v) (I(u + x, v + y) - I(u, v))^2$$

Diagram illustrating the components of the Harris corner detection formula:

- window**: Points to the summation indices  $u$  and  $v$ .
- Shifted intensities**: Points to the term  $I(u + x, v + y)$ .
- intensity**: Points to the term  $I(u, v)$ .



# Corner Detection: Harris

Shifted intensity is approximated using a Taylor Expansion:

$$I(u + x, v + y) \approx I(u, v) + I_x(u, v)x + I_y(u, v)y$$

So, at the end:

$$S(x, y) \approx \sum_u \sum_v w(u, v) (I_x(u, v)x + I_y(u, v)y)^2,$$

We can write this in matrix format as:

$$S(x, y) \approx \begin{pmatrix} x & y \end{pmatrix} A \begin{pmatrix} x \\ y \end{pmatrix},$$

, where A is the 2D structural tensor

$$A = \sum_u \sum_v w(u, v) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} \langle I_x^2 \rangle & \langle I_x I_y \rangle \\ \langle I_x I_y \rangle & \langle I_y^2 \rangle \end{bmatrix}$$

# Corner Detection: Harris

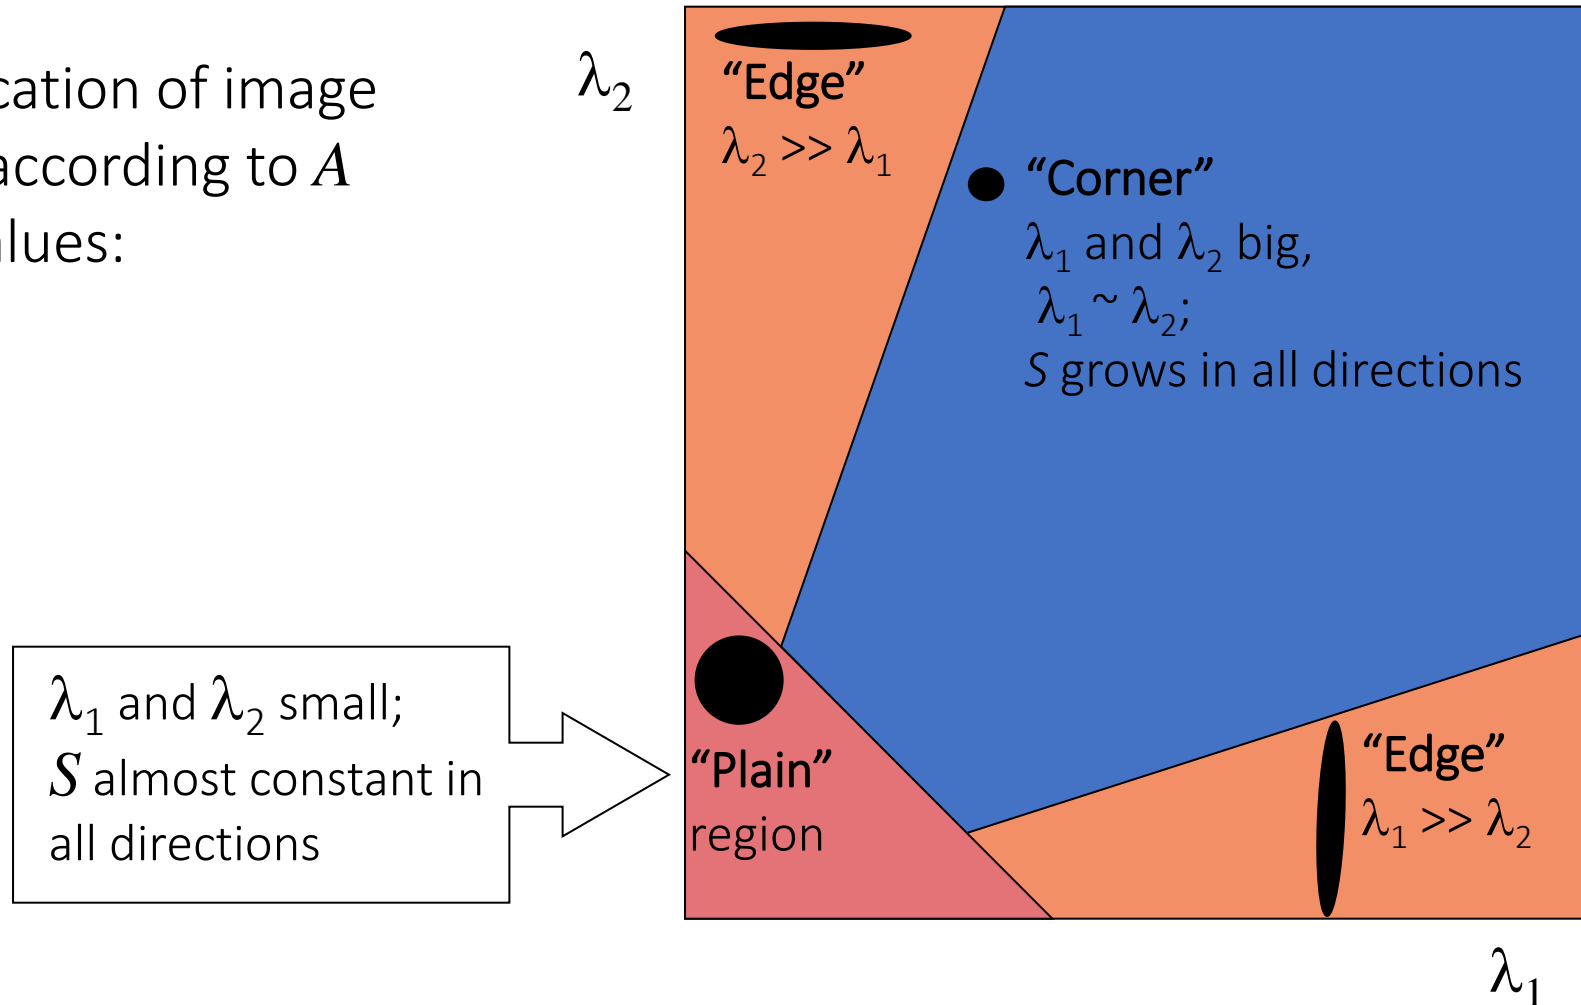
We change the problem of examining intensity changes due to translations to analyze the behaviour of matrix  $A \rightarrow$  analysis of eigenvalues

$\lambda_1, \lambda_2$  eigenvalues of  $A$



# Corner Detection: Harris

Classification of image points according to  $A$  eigenvalues:





# Corner Detection: Harris

Response function at corners ( $R$ ):

$$R = \det(A) - k (\text{trace } A)^2$$


$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

where  $k$  is a constant value (empiric)  $k = [0.04, 0.06]$



# Corner Detection: Harris

- $R$  depends only of  $A$  eigenvalues
- $R$  is big at corners
- $R$  is negative with high value at edges
- $|R|$  is small at plain regions



# Corner Detection: Harris

- First derivatives at an image point  $(u,v)$ :

$$I_x(u,v) = \frac{\partial I}{\partial x}(u,v)$$

$$I_y(u,v) = \frac{\partial I}{\partial y}(u,v)$$

- We can compute:

$$A(u,v) = I_x^2(u,v),$$

$$B(u,v) = I_x I_y(u,v),$$

$$C(u,v) = I_y^2(u,v)$$

- Local structure matrix ( $M$ )  
[a.k.a.  $A$ ]

$$M = \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} = \begin{pmatrix} A & C \\ C & B \end{pmatrix}$$

- Smoothing with a gaussian ( $G$ )

$$\bar{M} = \begin{pmatrix} A * G & C * G \\ C * G & B * G \end{pmatrix} = \begin{pmatrix} \bar{A} & \bar{C} \\ \bar{C} & \bar{B} \end{pmatrix}$$

# Corner Detection: Harris

- Diagonal of  $\overline{M}$

$$\overline{M} = \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- Where  $\lambda_1, \lambda_2$  are the eigenvalues of  $\overline{M}$  defined by:

$$\frac{1}{2} \left( \overline{A} + \overline{B} \pm \sqrt{\overline{A}^2 - 2\overline{A}\overline{B} + \overline{B}^2 + 4\overline{C}^2} \right)$$

- Describes a point according to eigenvalues, using corners response function

$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

- A good corner has big changes of intensity in all directions  $\rightarrow$  R should be big and positive.

# Corner Detection: Harris


Original






# Corner Detection: Harris

$R$



# Corner Detection: Harris

Points with  $R > \text{threshold}$





# Corner Detection: Harris

$R$  local maxima



# Corner Detection: Harris


Final result



# Corner Detection: Harris-Laplace


- Properties:

- Rotation invariant:



- Partial invariance to affine intensity changes (derivatives):

- Invariance to shifts in intensity  $I \rightarrow I+b$
    - Contrast change:  $I \rightarrow aI$



# Corner Detection: Harris-Laplace

- Combines Harris with a gaussian scale-space.
- We use gaussian Windows with predetermined scales
- We choose the scale that maximizes LoG in this range



- We obtain both the corners and the scale in which it is better represented



# Corner Detection: Harris-Affine

- Initial detection using Harris-Laplace
- Affine shape estimated using 2D structure matrix
- Normalize affine regions to a circular shape
- Detect new corner position and scales in the previous image
- If eigenvalues change, go back to point 2



# Harris Corner Detector in OpenCV

```
import numpy as np
import cv2 as cv
filename = 'chessboard.png'
img = cv.imread(filename)
gray = cv.cvtColor(img,cv.COLOR_BGR2GRAY)
gray = np.float32(gray)
dst = cv.cornerHarris(gray,2,3,0.04)
dst = cv.dilate(dst,None)
# Threshold for an optimal value, it may vary depending on the image.
img[dst>0.01*dst.max()]=[0,0,255]
cv.imshow('dst',img)
if cv.waitKey(0) & 0xff == 27:
    cv.destroyAllWindows()
```


3

# Edge Detection

# Edge Detection

Why do contours appear in images?

- Change in Depth
- Change in Orientation
- Change in Reflectance
- Change in Illumination






# Edge Detection

- Boundaries (edges)
  - Image regions where gradient magnitude has maximum value
- Valleys / Crests (ridges)
  - Curve that represents a local maxima or minima


- Models

- Step

Sudden  
step edge



Slanted  
step edge



Smooth  
step edge




Planar  
edge



Line edge



Roof edge



- Creast

- Valley



# Edge Detection

- Gradient
  - Vector that points in the direction of the highest change

$$\text{grad}(I) = \nabla(I) = \left( \frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right) = (I_x, I_y)$$

- We can calculate its magnitude and orientation



$$|\nabla| = \sqrt{I_x^2 + I_y^2}$$

$$\theta = \arctan(I_y / I_x)$$

- Boundaries are associated to high magnitude gradients

# Edge Detection

- Smoothing / Regularization
  - Allows us to decrease noise and control analysis scale
  - First derivative increases noise. We can smooth before derivating (regularization)
  - Smoothing can be done using a Gaussian with good properties (certain frequencies are not amplified)
    - We can also derivate the convolution with the derivative of the gaussian



# Edge Detection

- Algorithms
  - Differential gradient operator
    - Roberts
    - Sobel
    - Prewitt
  - Laplacian of Gaussian
  - Canny

# Edge Detection

$$Prewitt(im) = \left( im * \begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix}, im * \begin{pmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix} \right)$$


$$Sobel(im) = \left( im * \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}, im * \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix} \right)$$

$$Roberts(im) = \left( im * \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}, im * \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix} \right)$$


# Edge Detection

$$Sobel(im) = \sqrt{\left( im * \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} \right)^2 + \left( im * \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix} \right)^2}$$

Original



Sobel



Original



Sobel



# Edge Detection

- Laplacian

$$\textit{Laplacian}(I) = \Delta(I) = \nabla^2(I) = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$


- Numerical approximation

$$\textit{Laplacian}(im) = im * \begin{pmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$


- Laplacian's zero crossings provide us image boundaries

# Edge Detection


- Laplacian
  - Disadvantage: result is noisier



Original




Laplacian




Crossings

- Solution: smoothing with a gaussian



Original



LoG



Crossings

- Advantage: provides as result closed contours



# Edge Detection

- Canny Edge Detector:
  - We calculate the gradient with gaussian derivatives
  - We apply non maximum suppression
    - Selection of a single entity out of many overlapping ones
  - Join and binarize
    - We define upper and lower thresholds
    - We accept all contours above the lower threshold that are connected to other boundaries above the upper threshold

# Edge Detection

- Canny Edge Detector:


1.- Original



2.- Norm of the gradient



3.- Thresholding



4.- Thinning (non-maximum suppression)



# Edge Detection

- Canny Edge Detector:
  - Scale



original



low  $\sigma$



high  $\sigma$


- Choice  $\sigma$  depends of the desired behaviour
  - High  $\sigma$  detects high scale boundaries
  - Low  $\sigma$  detects low scale boundaries (noisier appearance)

# Edge Detection

- Grouping:
  - Primitive detection from boundary parts or a set of points
    - Hough Transform for lines (SLHT)
    - Hough Transform for circles (CHT)
    - Generalized Hough Transform (GHT)

# Edge Detection

- Hough Transform for lines
  - Transform points associated to a pattern within a parameter space where they can be represented in a compact shape
  - Example for lines  $y = ax + b$





- Solution:  $\text{lin} \rightarrow \rho = x \sin\theta + y \cos\theta$

# Edge Detection

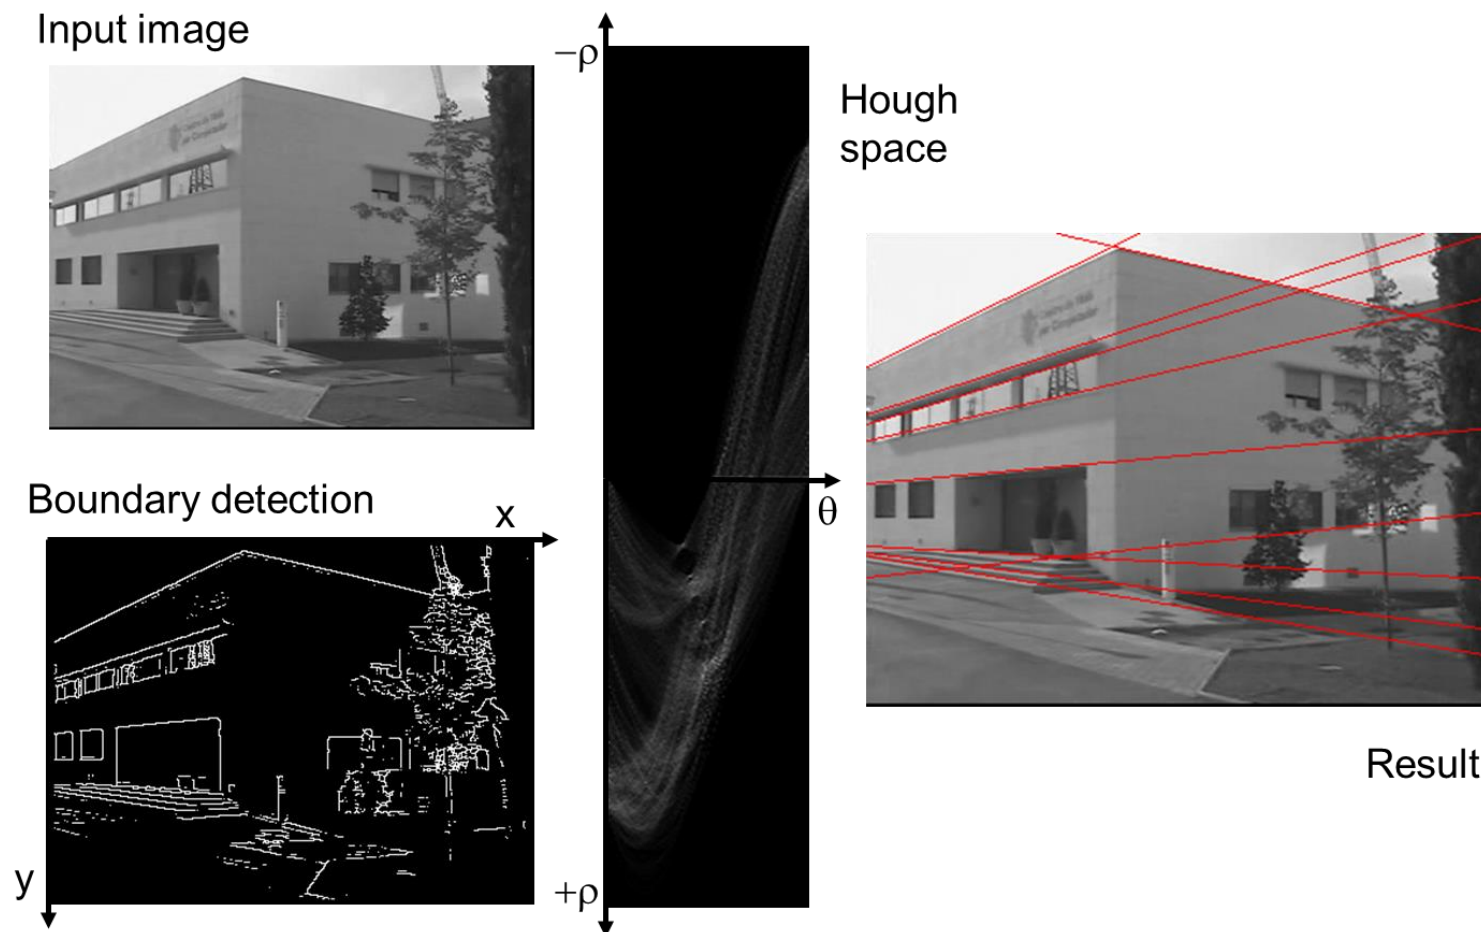
- Straight Line Hough Transform

- Key point detections: pixel selection according to local properties (gradient magnitude, orientation)
- Transformation mapping: each keypoint is mapped in the feature space (accumulation or voting array)
- Peak detection: local/global binarization in accumulation array



# Edge Detection

- Example





# Edge Detection

- Another Hough Transforms

- Circles (CHT):

- Tridimensional voting space  $(x,y,r)$
    - Each point contributes to this voting space within a cone

- General (GHT):

- Model definition:

- For an object (closed or open boundary) we define an inner center
      - For each boundary point we calculate the gradient (contour direction)
      - From the center to each point we calculate radii and angle
      - We store for each direction all radii-angle pairs

- Voting

- We generate image either of boundaries or from boundaries. We calculate gradients
      - For each point we vote all radii-angle associated to a particular direction

# Edge Detection

- Python implementation
  - Canny
    - `edge = cv2.Canny(image, low_th, high_th)`
  - Sobel
    - `edge = cv2.Sobel(image, precision_out_image, d_x, d_y)`
      - `d_x` and `d_y` specify if the first derivative of a specific direction is computed
  - Laplacian
    - `edge = cv.Laplacian(src_gray, precision_out_image, ksize)`
      - `ksize`: kernel size of the Sobel operator to be applied internally (commonly 3)
  - Hough Transform lines
    - `lines = cv.HoughLines(edges, rho, np.pi / 180, 150, None, 0, 0)`
      - `edges`: output of edge detector
      - `rho`: resolution of the parameter  $r$  in pixels (commonly 1)
      - `theta`: resolution of the parameter  $\theta$  in radians (commonly 1 degree,  $\pi/180$ )
      - `threshold`: minimum number of intersections to detect a line
      - `srn` and `stn`: set to 0

4

# Region Detectors

# Region Detection

- Laplacian of Gaussian
- Blob detection in binary images
- Maximally Stable Extremal Region (MSER)

# Region Detection



- Blobs
  - Regions in a digital image that differ in properties, such as brightness or color, compared to surrounding regions.
  - All the points in a blob can be considered in some sense to be similar to each other

# Region Detection

- Blobs
  - Two classes of blob detectors:
    - Differential: based on derivatives of the function with respect to position
    - Local-extrema based: finding the local maxima and minima of the function

# Region Detection


- Basic idea of Blob Detection: convolution of the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting *scale space*







# Region Detection

- Find maxima and minima of blob filter response in space and scale




# Region Detection: Laplacian of Gaussian

- Circularly symmetric operator for blob detection in 2D




$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Region Detection: Laplacian of Gaussian

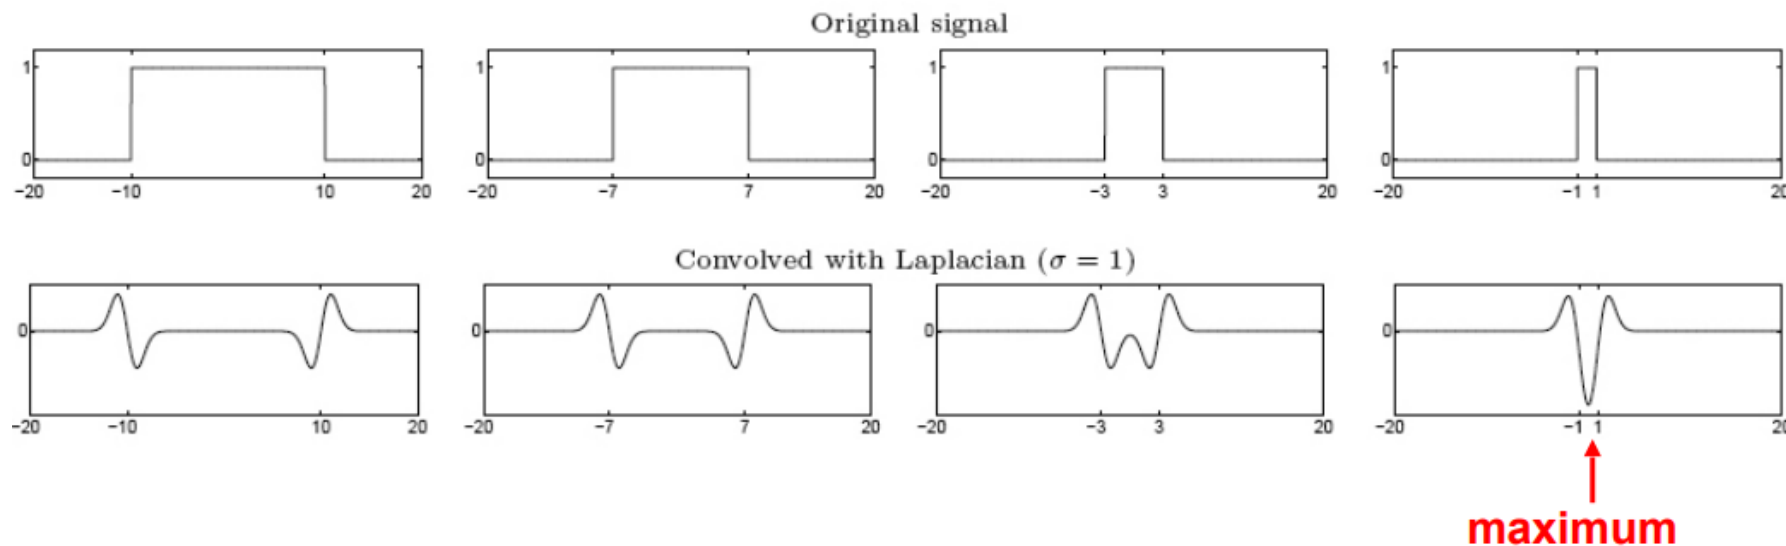


# Region Detection: Laplacian of Gaussian



# Region Detection: Laplacian of Gaussian

- Edge = ripple
- Blob = superposition of two ripples




**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian matches the scale of the blob

# Region Detection: Laplacian of Gaussian


**Scale selection:** we want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response

BUT Laplacian response decays as scale increases



# Region Detection: Laplacian of Gaussian


The response of a derivative of Gaussian filter to a perfect step Edge decreases as  $\sigma$  increases





# Region Detection: Laplacian of Gaussian

The response of a derivative of Gaussian filter to a perfect step Edge decreases as  $\sigma$  increases




How to make it scale-invariant?






- Multiply Gaussian derivative by  $\sigma$
- As Laplacian is the second Gaussian derivative, it should be multiplied by  $\sigma^2$

# Region Detection: Laplacian of Gaussian






Original signal



Unnormalized Laplacian response





Scale-normalized Laplacian response



maximum

# Region Detection: Laplacian of Gaussian

- Circularly symmetric operator for blob detection in 2D




Scale-normalized:


$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

# Region Detection: Laplacian of Gaussian


- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?



image



Laplacian




# Region Detection: Laplacian of Gaussian

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- For maximum response: align the zeros of the Laplacian with the circle
- The Laplacian in 2-D is given by (up to scale):


$$(x^2 + y^2 - 2\sigma^2)e^{-(x^2 + y^2)/(2\sigma^2)}$$

- Therefore, the maximum response occurs at  $\sigma = r / \sqrt{2}$ .



# Region Detection: Laplacian of Gaussian

- We define the characteristic scale of a blob as the scale that produces peak of Laplacian response in the blob center





# Region Detection: Laplacian of Gaussian

- 1/ Convolve image with scale-normalized Laplacian at several scales




sigma = 11.9912




# Region Detection: Laplacian of Gaussian

- 2/ Find maxima of squared Laplacian response in scale-space




# Region Detection: Laplacian of Gaussian




# Region Detection: Blob Detection in Binary Images

- Blob in binary images:
  - Collection of white pixels connected to each other
  - The same label is given to all the pixels belonging to the same blob
  - Useful to count the number of object in the scene




Original image



Labeled image


# Region Detection: Blob Detection in Binary Images

- One or two objects?




# Region Detection: Blob Detection in Binary Images

- Connected component analysis



1


Connectivity = 4



Each pixel has 4  
possible neighbors

2

Connectivity = 8





Each pixel has 8  
possible neighbors

# Region Detection: Blob Detection in Binary Images

- Connected component analysis algorithm: First stage
  1. Create an image with the same dimension and initialized to 0
  2. Label = 1
  3. Move over the image pixel by pixel in direction left-right and top-down
  4. if  $In(x,y) == 0$  go to the next pixel
  5. else check the neighbor's labels in Out(x,y) image
    - i. if all the neighbours are 0 assign Label and Label += 1
    - ii. elif only one label in the neighbours assign this label to Out(x,y)
    - iii. else there are more than one label assign the lowest one and add to the equivalent table the neighbor's labels



# Region Detection: Blob Detection in Binary Images






# Region Detection: MSER

## **Maximally Stable Extremal Regions**

- MSER regions are connected areas characterized by almost uniform intensity, surrounded by contrasting background.
- They are constructed through a process of trying multiple thresholds
- The selected regions are those that maintain unchanged shapes over a large set of thresholds.

# Region Detection: MSER

## Examples of MSER regions




# Region Detection: MSER

## How to create MSER regions



intensity image




shown as a surface function


Do you remember watersheds?

# Region Detection: MSER

## How to create MSER regions



Threshold simulation



External regions (we store region are per each threshold value)

# Region Detection: MSER

## **How to create MSER regions**

- For each threshold, compute the connected binary regions
- Calculate Area at each threshold value
- Analyze this function to determine those regions that have a similar value over multiple thresholds
- Regions detected at different thresholds have different areas

# Region Detection: Python Example (LoG)

```
import numpy as np
import matplotlib.pyplot as plt
from skimage.feature import blob_log
import cv2 as cv2

im = cv2.imread('C:/Users/Jorge/Desktop/FCV/almendras.png',cv2.IMREAD_GRAYSCALE)
cv2.imshow('orig',im)
cv2.waitKey(0)

#binarize the image
im_bw, th = cv2.threshold(im,200,255,cv2.THRESH_BINARY_INV)
print(th.shape)
cv2.imshow('bin',th)
cv2.waitKey(0)
```



# Region Detection: Python Example (LoG)

```
blobs = blob_log(th, max_sigma=30, min_sigma = 3, num_sigma=2, threshold=0.3, overlap = 0.1)
fig, ax = plt.subplots()
ax.imshow(th, cmap='gray')
for blob in blobs:
    y, x, area = blob
    ax.add_patch(plt.Circle((x, y), area*np.sqrt(2), color='r',
                           fill=False))

cv2.destroyAllWindows()
```





# Region Detection: Python Example (MSER)

```
detector = cv2.MSER_create()
keypoints = detector.detect(th, None)

# We draw all keypoints detected in the image
for keypoint in keypoints:
    radius = int(0.5 * keypoint.size)
    x, y = np.int64(keypoint.pt)
    cv2.circle(im, (x, y), radius, (0, 255, 255), 2)

# We show the original image and the one with the regions marked
cv2.imshow('Imágenes', np.hstack([copia, im]))
```



# Fundamentals of Computer Vision

Unit 6: Feature Extraction

Jorge Bernal