

Fundamentals of Computer Vision

Unit 6: Feature Extraction

Jorge Bernal

Index

01

1. Introduction

02

2. Corner
detectors

03

3. Edge
detectors

04

4. Blob
detectors

1

Introduction

Introduction

Definition of feature:

- Piece of information that is useful to solve a given task
- Interesting part of the image



Types of features:

- **Global:** global properties of the whole image
 - Mean grey level, mean colour, main colours, histogram
- **Local:** properties of a part of the image with their own entity
 - Points, edges, regions

Introduction

GLOBAL

LOCAL



Introduction

- **Local features:**
 - Part of an image that differs from its surroundings.
 - They are associated to a change in a certain property (intensity, color, texture)
 - Examples:
 - Points (corners, interest points)
 - Edges, ridges
 - Small regions (blobs)

Introduction

- How can we find local features?
 - Feature detection/extraction algorithms
- Detection/Extraction: locate the position of the feature
- Description (Unit 7): measures that are taken from the detected feature that allow us to distinguish it or compare with others

Introduction

- Why do we use features?
 - They have been used with success in several disciplines and applications:
 - Edge detection associated to roads in aerial images
 - Quality control
 - Polyp Detection
 - Interest points play a key role for certain Applications:
 - Tracking
 - 3D reconstruction
 - They are a first step to achieve a robust image representation:
 - Object recognition
 - Scene classification
 - Texture analysis
 - Image search

Introduction

- Ideal properties:
 - Repetability:
 - Invariance to transformations
 - Robustness
 - Differentiation (highly different from another)
 - Precise localization
 - Enough points for the needed task
 - Efficient
- Scale: very important factor to achieve robustness, invariance and precision.
Allows us to work with different images at several distances.


Introduction

Feature Detector				Rotation invariant	Scale invariant	Affine invariant	Localization			
	Corner	Blob	Region				Repeatability	accuracy	Robustness	Efficiency
Harris	✓			✓			+++	+++	+++	++
Hessian		✓		✓			++	++	++	+
SUSAN	✓			✓			++	++	++	++++
Harris-Laplace	✓	(✓)		✓	✓		+++	+++	++	+
Hessian-Laplace	(✓)	✓		✓	✓		+++	+++	+++	+
DoG	(✓)	✓		✓	✓		++	++	++	++
SURF	(✓)	✓		✓	✓		++	++	++	++++
Harris-Affine	✓	(✓)		✓	✓	✓	+++	+++	++	++
Hessian-Affine	(✓)	✓		✓	✓	✓	+++	+++	+++	++
Salient Regions	(✓)	✓		✓	✓	(✓)	+	+	++	+
Edge-based	✓			✓	✓	✓	+++	+++	+	+
MSER			✓	✓	✓	✓	+++	+++	++	++++
Intensity-based			✓	✓	✓	✓	++	++	++	++
Superpixels			✓	✓	(✓)	(✓)	+	+	+	+

2


Corner Detection

Corner Detection




Plain region

No changes in all directions



Edge

No change in edge direction



Corner

Significant changes in all directions

Corner Detection

- Harris (1988): Based on the analysis of the 2D structural tensor (second derivative matrix, second moment matrix)
- SUSAN (Smallest Univalued Segment Assimilating Nucleus): morphologic focus
- Harris-Laplace: Use of Harris for a first detection; scale is fixed using laplacian
- Harris-Affine: Use of Harris-Laplace; then it tries to estimate the most affine shape (with an ellipse that is later normalized to a circle)

Corner Detection: Harris


- In an image intensity corner, intensity changes significantly in all directions.
- Here we are focused in intensity changes in a local window.
- We use SSD: sum of squared differences

$$S(x, y) = \sum_u \sum_v w(u, v) (I(u + x, v + y) - I(u, v))^2$$


Diagram illustrating the components of the Harris corner detection formula:

- window**: Points to the summation indices u and v .
- Shifted intensities**: Points to the term $I(u + x, v + y)$.
- intensity**: Points to the term $I(u, v)$.

Window function $w(u, v) =$



1 inside, 0 outside



gaussian

Corner Detection: Harris

Shifted intensity is approximated using a Taylor Expansion:

$$I(u+x, v+y) \approx I(u, v) + I_x(u, v)x + I_y(u, v)y$$

So, at the end:

$$S(x, y) \approx \sum_u \sum_v w(u, v) (I_x(u, v)x + I_y(u, v)y)^2,$$

We can write this in matrix format as:

$$S(x, y) \approx \begin{pmatrix} x & y \end{pmatrix} A \begin{pmatrix} x \\ y \end{pmatrix},$$

, where A is the 2D structural tensor

$$A = \sum_u \sum_v w(u, v) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} \langle I_x^2 \rangle & \langle I_x I_y \rangle \\ \langle I_x I_y \rangle & \langle I_y^2 \rangle \end{bmatrix}$$

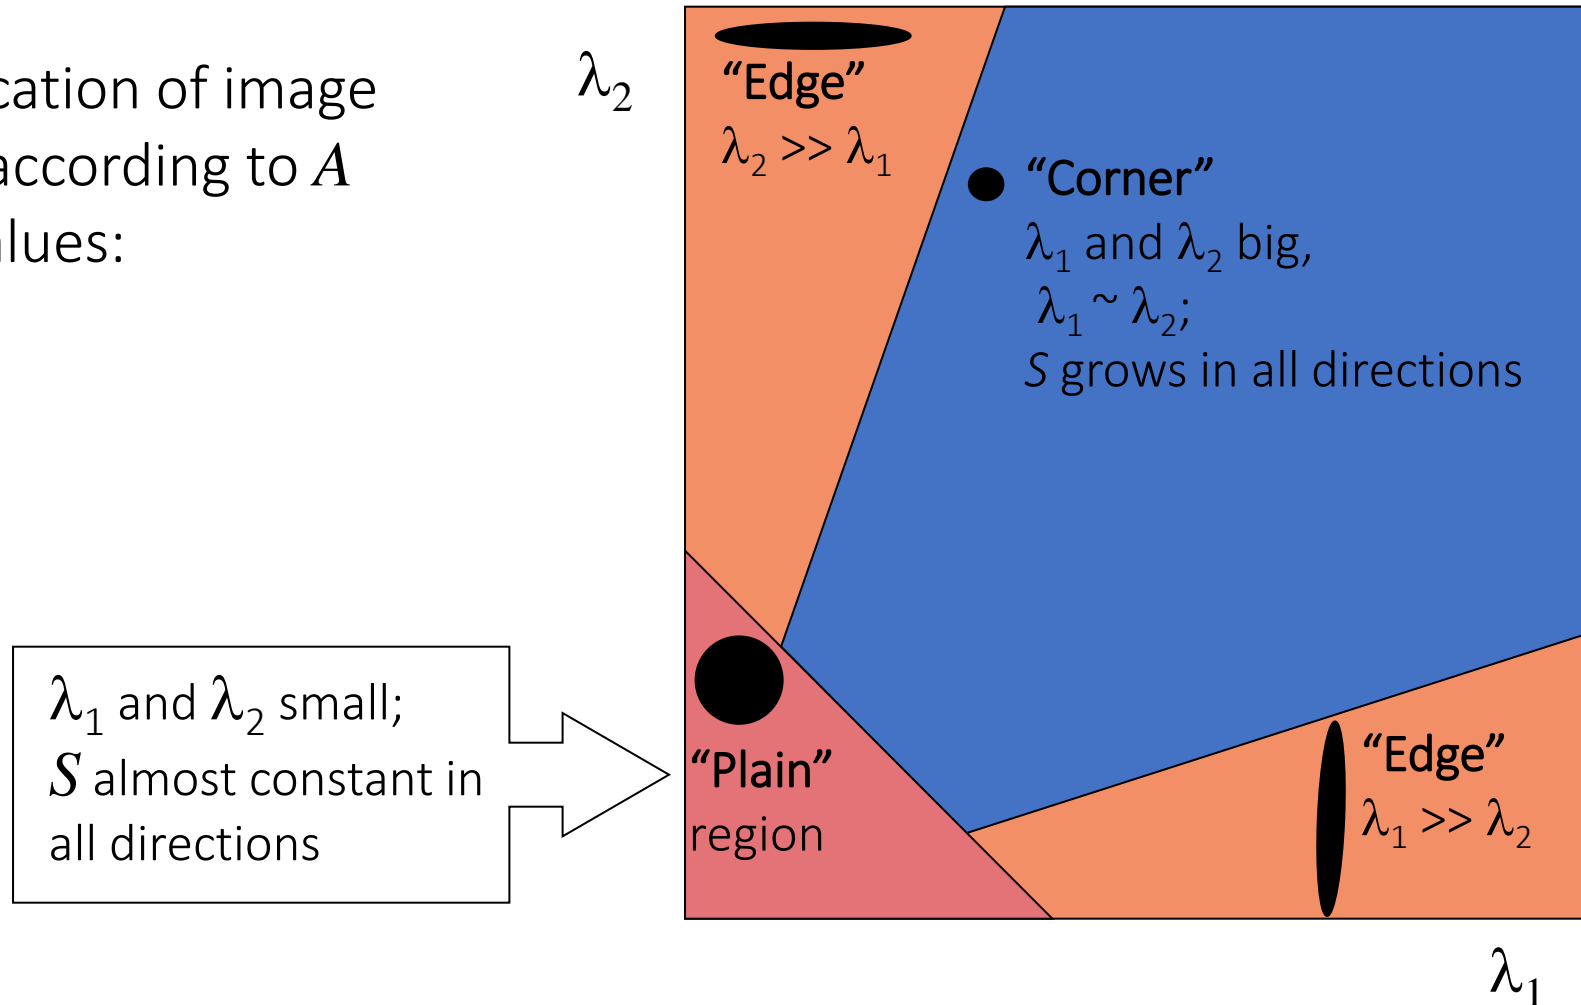
Corner Detection: Harris

We change the problem of examining intensity changes due to translations to analyze the behaviour of matrix $A \rightarrow$ analysis of eigenvalues

λ_1, λ_2 eigenvalues of A

Corner Detection: Harris

Classification of image points according to A eigenvalues:





Corner Detection: Harris

Response function at corners (R):

$$R = \det(A) - k (\text{trace } A)^2$$


$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

where k is a constant value (empiric) $k = [0.04, 0.06]$



Corner Detection: Harris

- R depends only of A eigenvalues
- R is big at corners
- R is negative with high value at edges
- $|R|$ is small at plain regions



Corner Detection: Harris

- First derivatives at an image point (u,v) :

$$I_x(u,v) = \frac{\partial I}{\partial x}(u,v)$$

$$I_y(u,v) = \frac{\partial I}{\partial y}(u,v)$$

- We can compute:

$$A(u,v) = I_x^2(u,v),$$

$$B(u,v) = I_x I_y(u,v),$$

$$C(u,v) = I_y^2(u,v)$$

- Local structure matrix (M)
[a.k.a. A]

$$M = \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} = \begin{pmatrix} A & C \\ C & B \end{pmatrix}$$

- Smoothing with a gaussian (G)

$$\bar{M} = \begin{pmatrix} A * G & C * G \\ C * G & B * G \end{pmatrix} = \begin{pmatrix} \bar{A} & \bar{C} \\ \bar{C} & \bar{B} \end{pmatrix}$$

Corner Detection: Harris

- Diagonal of \overline{M}

$$\overline{M} = \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- Where λ_1, λ_2 are the eigenvalues of \overline{M} defined by:

$$\frac{1}{2} \left(\overline{A} + \overline{B} \pm \sqrt{\overline{A}^2 - 2\overline{A}\overline{B} + \overline{B}^2 + 4\overline{C}^2} \right)$$

- Describes a point according to eigenvalues, using corners response function

$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

- A good corner has big changes of intensity in all directions \rightarrow R should be big and positive.


Corner Detection: Harris

Original




Corner Detection: Harris

R



Corner Detection: Harris

Points with $R > \text{threshold}$



Corner Detection: Harris

R local maxima



Corner Detection: Harris


Final result



Corner Detection: Harris-Laplace


- Properties:

- Rotation invariant:



- Partial invariance to affine intensity changes (derivatives):

- Invariance to shifts in intensity $I \rightarrow I+b$
 - Contrast change: $I \rightarrow aI$



Corner Detection: Harris-Laplace

- Combines Harris with a gaussian scale-space.
- We use gaussian Windows with predetermined scales
- We choose the scale that maximizes LoG in this range



- We obtain both the corners and the scale in which it is better represented

Corner Detection: Harris-Affine

- Initial detection using Harris-Laplace
- Affine shape estimated using 2D structure matrix
- Normalize affine regions to a circular shape
- Detect new corner position and scales in the previous image
- If eigenvalues change, go back to point 2



Harris Corner Detector in OpenCV

```
import numpy as np
import cv2 as cv
filename = 'chessboard.png'
img = cv.imread(filename)
gray = cv.cvtColor(img,cv.COLOR_BGR2GRAY)
gray = np.float32(gray)
dst = cv.cornerHarris(gray,2,3,0.04)
dst = cv.dilate(dst,None)
# Threshold for an optimal value, it may vary depending on the image.
img[dst>0.01*dst.max()]=[0,0,255]
cv.imshow('dst',img)
if cv.waitKey(0) & 0xff == 27:
    cv.destroyAllWindows()
```

3

Edge Detection

Edge Detection

Why do contours appear in images?

- Change in Depth
- Change in Orientation
- Change in Reflectance
- Change in Illumination




Edge Detection

- Boundaries (edges)
 - Image regions where gradient magnitude has maximum value
- Valleys / Crests (ridges)
 - Curve that represents a local maxima or minima


- Models

- Step

Sudden
step edge



Slanted
step edge



Smooth
step edge




Planar
edge



Line edge



Roof edge



- Creast

- Valley



Edge Detection

- Gradient
 - Vector that points in the direction of the highest change

$$\text{grad}(I) = \nabla(I) = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right) = (I_x, I_y)$$

- We can calculate its magnitude and orientation



$$|\nabla| = \sqrt{I_x^2 + I_y^2}$$

$$\theta = \arctan(I_y / I_x)$$

- Boundaries are associated to high magnitude gradients

Edge Detection

- Smoothing / Regularization
 - Allows us to decrease noise and control analysis scale
 - First derivative increases noise. We can smooth before derivating (regularization)
 - Smoothing can be done using a Gaussian with good properties (certain frequencies are not amplified)
 - We can also derivate the convolution with the derivative of the gaussian



Edge Detection

- Algorithms
 - Differential gradient operator
 - Roberts
 - Sobel
 - Prewitt
 - Laplacian of Gaussian
 - Canny

Edge Detection

$$Prewitt(im) = \left(im * \begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix}, im * \begin{pmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix} \right)$$


$$Sobel(im) = \left(im * \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}, im * \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix} \right)$$

$$Roberts(im) = \left(im * \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}, im * \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix} \right)$$


Edge Detection

$$Sobel(im) = \sqrt{\left(im * \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} \right)^2 + \left(im * \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix} \right)^2}$$

Original



Sobel



Original



Sobel



Edge Detection

- Laplacian

$$\textit{Laplacian}(I) = \Delta(I) = \nabla^2(I) = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$


- Numerical approximation

$$\textit{Laplacian}(im) = im * \begin{pmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$


- Laplacian's zero crossings provide us image boundaries

Edge Detection


- Laplacian
 - Disadvantage: result is noisier



Original




Laplacian




Crossings

- Solution: smoothing with a gaussian



Original



LoG



Crossings

- Advantage: provides as result closed contours

Edge Detection

- Canny Edge Detector:
 - We calculate the gradient with gaussian derivatives
 - We apply non maximum suppression
 - Selection of a single entity out of many overlapping ones
 - Join and binarize
 - We define upper and lower thresholds
 - We accept all contours above the lower threshold that are connected to other boundaries above the upper threshold

Edge Detection

- Canny Edge Detector:


1.- Original



2.- Norm of the gradient



3.- Thresholding



4.- Thinning (non-maximum suppression)



Edge Detection

- Canny Edge Detector:
 - Scale



original



low σ



high σ


- Choice σ depends of the desired behaviour
 - High σ detects high scale boundaries
 - Low σ detects low scale boundaries (noisier appearance)

Edge Detection

- Grouping:
 - Primitive detection from boundary parts or a set of points
 - Hough Transform for lines (SLHT)
 - Hough Transform for circles (CHT)
 - Generalized Hough Transform (GHT)

Edge Detection

- Hough Transform for lines
 - Transform points associated to a pattern within a parameter space where they can be represented in a compact shape
 - Example for lines $y = ax + b$





- Solution: $\text{lin} \rightarrow \rho = x \sin\theta + y \cos\theta$

Edge Detection

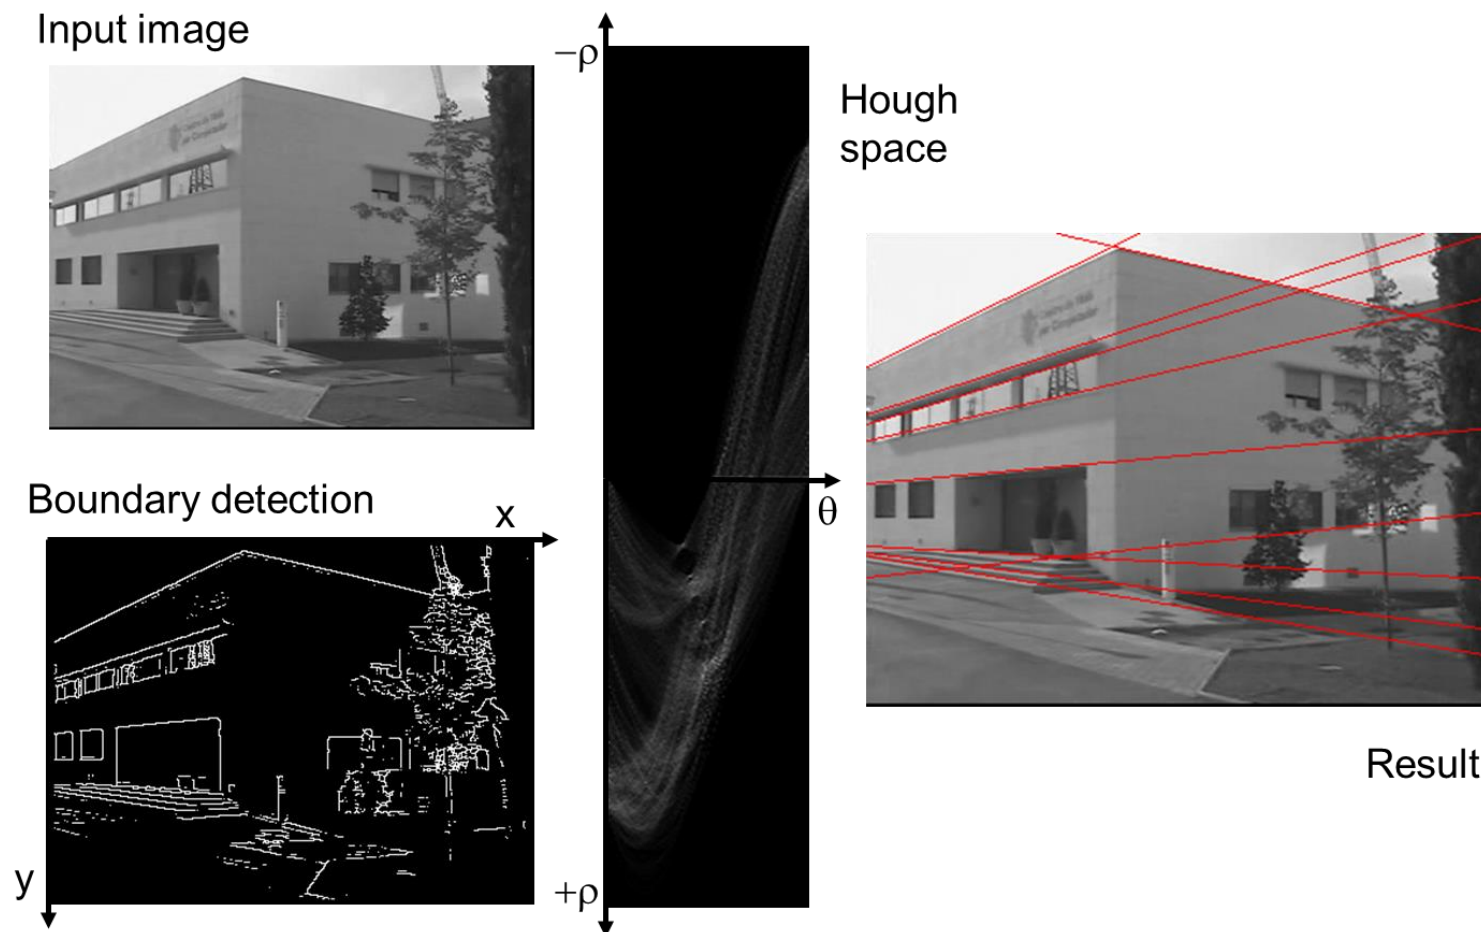
- Straight Line Hough Transform

- Key point detections: pixel selection according to local properties (gradient magnitude, orientation)
- Transformation mapping: each keypoint is mapped in the feature space (accumulation or voting array)
- Peak detection: local/global binarization in accumulation array



Edge Detection

- Example



Edge Detection

- Another Hough Transforms

- Circles (CHT):

- Tridimensional voting space (x,y,r)
 - Each point contributes to this voting space within a cone

- General (GHT):

- Model definition:

- For an object (closed or open boundary) we define an inner center
 - For each boundary point we calculate the gradient (contour direction)
 - From the center to each point we calculate radii and angle
 - We store for each direction all radii-angle pairs

- Voting

- We generate image either of boundaries or from boundaries. We calculate gradients
 - For each point we vote all radii-angle associated to a particular direction

Edge Detection

- Python implementation
 - Canny
 - `edge = cv2.Canny(image, low_th, high_th)`
 - Sobel
 - `edge = cv2.Sobel(image, precision_out_image, d_x, d_y)`
 - `d_x` and `d_y` specify if the first derivative of a specific direction is computed
 - Laplacian
 - `edge = cv.Laplacian(src_gray, precision_out_image, ksize)`
 - `ksize`: kernel size of the Sobel operator to be applied internally (commonly 3)
 - Hough Transform lines
 - `lines = cv.HoughLines(edges, rho, np.pi / 180, 150, None, 0, 0)`
 - `edges`: output of edge detector
 - `rho`: resolution of the parameter r in pixels (commonly 1)
 - `theta`: resolution of the parameter θ in radians (commonly 1 degree, $\pi/180$)
 - `threshold`: minimum number of intersections to detect a line
 - `srn` and `stn`: set to 0

Fundamentals of Computer Vision

Unit 6: Feature Extraction

Jorge Bernal